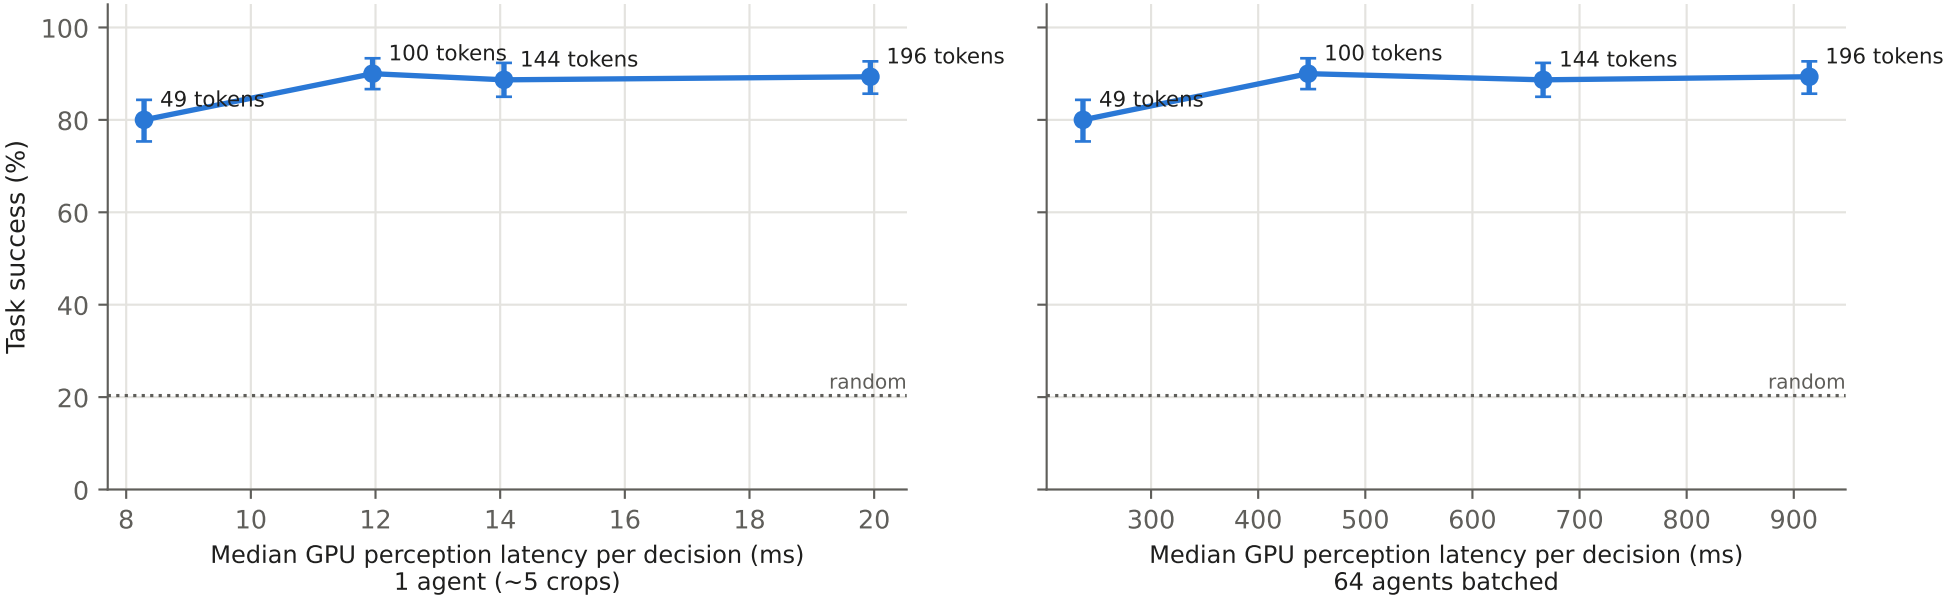
Efficiency vs task success · Tesla T4 · float16 · sdpa · preprocessed crops on GPU → logits
Task success (%)
100
80
60
40
20
0
49 tokens
100 tokens
144 tokens
196 tokens
random
8
10
12
14
16
18
20
Median GPU perception latency per decision (ms)
1 agent (~5 crops)
49 tokens
100 tokens
144 tokens
196 tokens
random
300
400
500
600
700
800
900
Median GPU perception latency per decision (ms)
64 agents batched